Exercises for the course
Deep Learning 2
Summer Semester 2025

Machine Learning Group
Faculty IV – Electrical Engineering and Computer Science
Technische Universität Berlin

Exercise Sheet **Representation Learning**

Exercise 1: Contrastive Loss (20 P)

Given the SimCLR loss from the lecture for all views i, j from the same samples in a minibatch (MB).

$$\mathcal{L} = -\frac{1}{N} \sum_{i,j \in MB} \log \frac{\exp\left(\text{sim}\left(\mathbf{z}_i, \mathbf{z}_j\right)/\tau\right)}{\sum_{k=1}^{2N} \mathbb{1}_{[k \neq i]} \exp\left(\text{sim}\left(\mathbf{z}_i, \mathbf{z}_k\right)/\tau\right)} \tag{1}$$

with $\text{sim}(u, v) = \frac{u^T v}{\|u\|\|v\|}$ being the cosine similarity τ a scalar and N the number of samples.

a) Rewrite the loss explicitly into the following form:

$$\tau\mathcal{L} = \mathcal{L}_a + \mathcal{L}_d$$

with $\mathcal{L}_a = -\frac{1}{N} \sum_{i,j \in MB} \text{sim}(\mathbf{z}_i, \mathbf{z}_j)$.

What is the purpose of $\mathcal{L}_a$ and $\mathcal{L}_d$ in the loss?

$$\mathcal{L}_d = \frac{\tau}{N} \sum_{i} \log \sum_{k=1}^{2N} \mathbb{1}_{[k \neq i]} \exp\left(\text{sim}\left(\mathbf{z}_i, \mathbf{z}_k\right)/\tau\right) \tag{2}$$

$\mathcal{L}_a$ encourages representations of augmented views to be consistent (alignment) while $\mathcal{L}_d$ encourages representations (or a random subset of them) to match a prior distribution (of high entropy). $\mathcal{L}_d$ prevents a representation collapse and helps to distribute the representations across the whole space (uniformity).

b) How does the parameter τ influences the distance between representations?

τ is the temperature parameter that is usually used to calibrate a softmax distribution. In contrast to regular classification, the logits are in this case determined by the cosine similarities and therefore bounded in the range $[-1, 1]$. Using a low temperature makes the distribution to have less entropy. Therefore, the differences in the probabilities (distances/similarities) after the softmax also become larger. Using a larger temperature leads to a more uniform distribution of the similarities/distances after the softmax.

Exercise 2: Lecture Questions (20 P)

a) What is a pretext task? Give four examples for pretext tasks.

A pretext task is an auxiliary task that has an objective that does not require explicit labels. The pretext task is usually performed on a property that is inherent in the dataset itself. Examples:

- Masking and Reconstruction
- Contrastive Learning Objective
- Color prediction (grayscale to color)
- Jigsaw Puzzle task
- Rotation prediction
- Super-resolution

b) What is a representation collapse and how is it prevented in SimCLR?

A representation collapse means that the encoder maps all representations to a single point (e.g $0 \in \mathbb{R}^d$). SimCLR uses negative samples which are used in the contrastive loss to minimize similarity (push representations apart) in order to avoid a representation collapse.

c) Given an image/text model with image encoder f and text encoder g which both produce a representation $z \in \mathbb{R}^d$, we want to perform zero-shot classification. Given text labels $t_1, \ldots, t_k$ that describe k classes and an image x, how do you compute the predicted class c?

$$c = \underset{i=1,\ldots,k}{\operatorname{argmax}} \frac{f(x)^T \cdot g(t_i)}{\|f(x)\| \|g(t_i)\|}$$

d) Name two other applications for representations from a pretext task other than using them for a classification downstream task.

- Clustering
- Semantic Search/retrieval tasks
- Anomaly Detection
- Finding independent components
- Matching them to other modalities

Exercise 3: Programming (60 P)

Download the programming files on ISIS and follow the instructions.